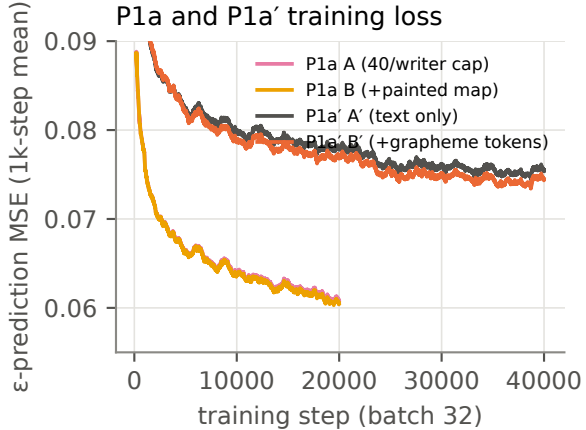
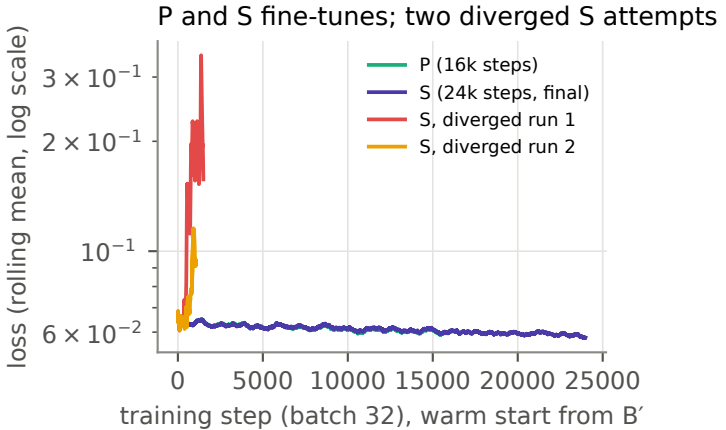

P1a and P1a′ training loss
P1a A (40/writer cap)
P1a B (+painted map)
P1a′ A′ (text only)
P1a′ B′ (+grapheme tokens)
ε-prediction MSE (1k-step mean)
0.09
0.08
0.07
0.06
0
10000
20000
30000
40000
training step (batch 32)


P and S fine-tunes; two diverged S attempts
P (16k steps)
S (24k steps, final)
S, diverged run 1
S, diverged run 2
loss (rolling mean, log scale)
3 × 10⁻¹
2 × 10⁻¹
10⁻¹
6 × 10⁻²
0
5000
10000
15000
20000
25000
training step (batch 32), warm start from B′